

Report

Krish Dange

November 2025

1 Summary of Models Used & System Architecture

1.1 PCA Preprocessing

- MNIST images have 784 raw pixel features.
- PCA was used to reduce dimensionality.
- Reduced to 100 principal components.
- Shapes after PCA:
 - Train: (10002, 100)
 - Validation: (2499, 100)
- Total variance preserved: 91.56%.
- PCA greatly reduced computation time and improved model stability.

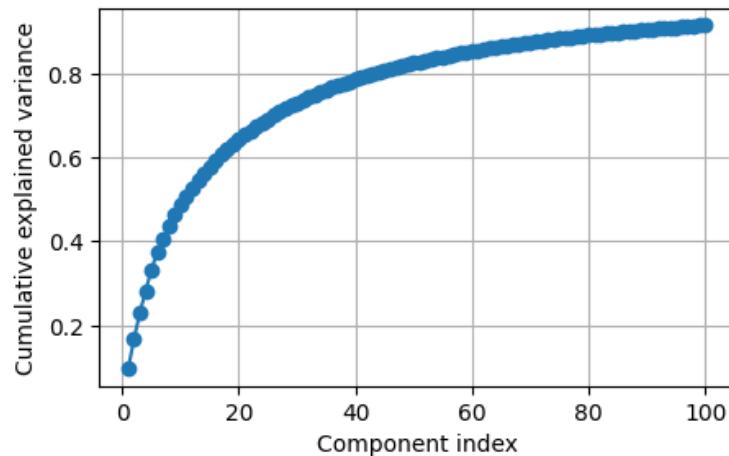


Figure 1: Cumulative explained variance for 100 PCA components

1.2 Models Implemented

Linear Models

- Softmax Logistic Regression
- Linear SVM (One-vs-Rest)

Non-linear Models

- Random Forest
- One-vs-Rest Gradient Boosting

1.3 Initial Model Performance

Model	Weighted F1	Accuracy	Train Time (s)
Logistic Regression	0.8561	0.8575	0.35
Random Forest	0.7071	0.7135	97.43
One-vs-Rest GB	0.5753	0.5722	19.52
Linear SVM	0.8645	0.8667	0.72

Observation:

- Linear models clearly outperform scratch non-linear models.
- Thus, only Logistic Regression and SVM were tuned further.

2 Hyperparameter Tuning Summary

2.1 Softmax Logistic Regression

2.1.1 Learning Rate

Learning Rate	Weighted F1
0.02	0.9032
0.01	0.8944
0.005	0.8898

Best: **0.02**

2.1.2 Regularization

reg	Weighted F1
1e-4	0.9041
1e-5	0.9053
1e-6	0.9044

Best: **1e-5**

2.1.3 Batch Size

Batch Size	Weighted F1
128	0.9080
256	0.9032
512	0.8964

Notes:

- Batch 128 gives highest score but unstable.
- Batch 256 chosen for stability and consistent convergence.

Final Tuned LR Hyperparameters

- lr = 0.02
- reg = 1e-5
- batch_size = 256
- epochs = 400
- patience = 30
- Final Weighted F1: **0.905**

2.2 Linear SVM (One-vs-Rest)

2.2.1 C Tuning

C	Weighted F1
10	0.8627
50	0.8939
100	0.8980
300	0.9000
500	0.8988

Best: **C = 300**. Higher C improves fit until 300, then drops (overfitting).

2.2.2 Learning Rate

Learning Rate	Weighted F1
0.01	0.8964
0.02	0.9000
0.03	0.8988

Best: **0.02**

Final Tuned SVM Hyperparameters

- C = 300
- lr = 0.02
- batch_size = 256
- epochs = 400
- patience = 30
- Final Weighted F1: **0.900**

3 System Optimization Steps & Runtime Performance

- Applied PCA(100) to reduce dimensionality and speed up all models.
- Used efficient NumPy vectorized operations for all model training.
- Mini-batch gradient descent (batch = 256) for stable and fast convergence.
- Learning rate decay: $lr_t = \frac{lr}{1+0.005t}$.
- Early stopping (patience = 30) to prevent unnecessary epochs.
- All models trained under single-core CPU as required.

Final Validation Results

- Logistic Regression: **F1 = 0.905**, Acc = 0.906
- Linear SVM: **F1 = 0.900**, Acc = 0.899

4 Thoughts & Observations

- PCA drastically improved speed and stability; 100 components preserved 91.56% variance.
- After PCA, MNIST became more linearly separable; linear models outperformed non-linear ones.
- Logistic Regression and SVM responded strongly to learning rate and regularization tuning.
- Non-linear scratch models were slower and less accurate, showing the difficulty of manual tree-based implementations.
- Overall best pipeline: **PCA(100) + Softmax Logistic Regression**.
- Training times were extremely low (1–2 seconds), far below the 5-minute constraint.